

An Analysis of Multiprocessor Quicksort

Presented By:

Kevin Ellis, Braden Harrelson, Ryan Luig

What is a Serial and Parallel Program?

- A serial program utilizes one processor to accomplish its tasks
- A parallel program utilizes multiple processors in order to complete multiple tasks simultaneously

Project Goals

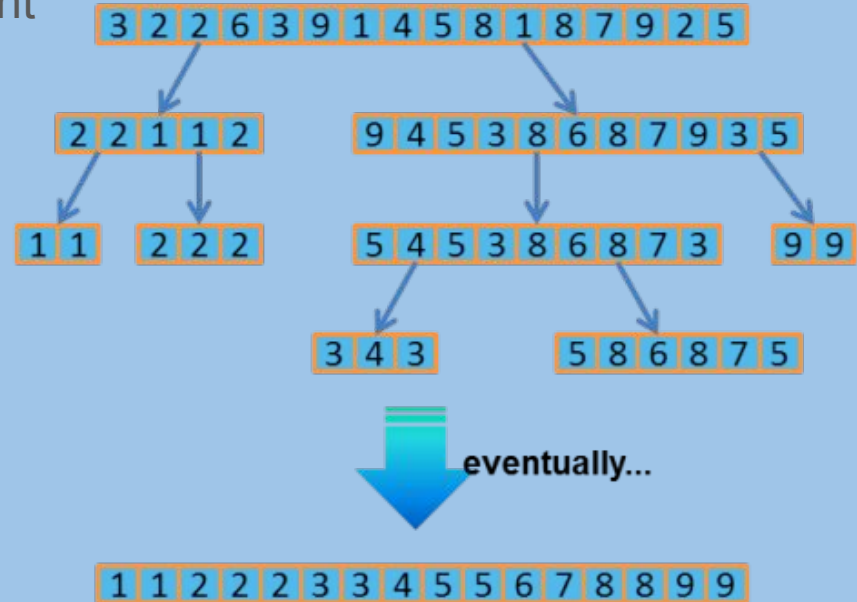
- Businesses store large amounts of data
 - client records, transaction list, etc.
- Develop sorting algorithm using MPI (Parallel library)
- Increase speed while being weary of memory usage
- Compare speed of serial vs parallel

Quicksort Algorithm vs Other Algorithms

- Quicksort is in place
 - less memory usage
- Complexity
 - Average complexity of $O(n \log n)$
 - worst case $O(n^2)$

How Serial Quicksort Works

- Select a pivot (We select the leftmost element)
- If less move to the left, more to the right
- Recursively call on left and right sides



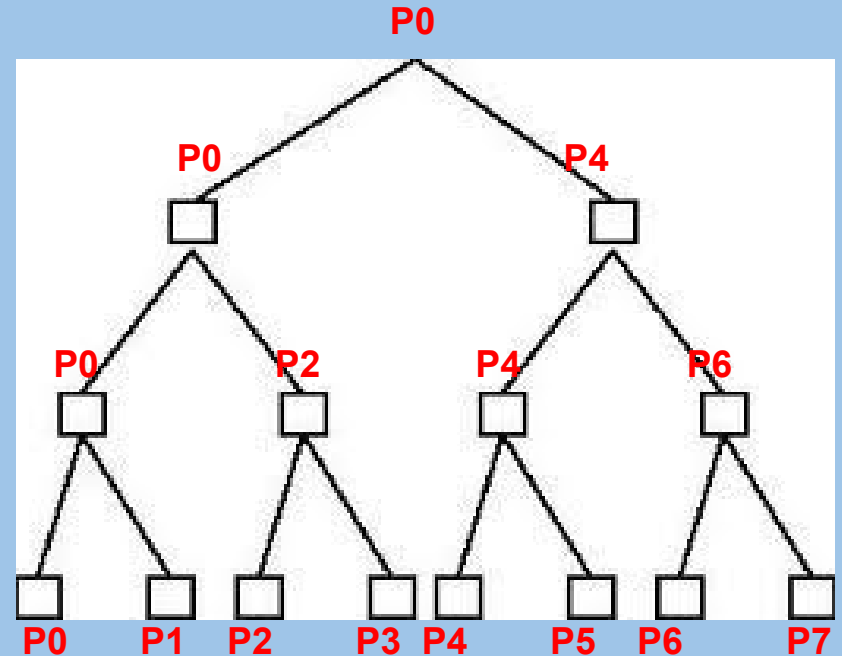
(Pic source: hpcwire.com)

Initial Approaches

- Sample sort
 - Issues with sentinel values
- Recursive pivot moving
 - Communication issues

Parallel Approach

- Divide and Conquer
- Sort Partition and Send
- Receive and Merge
- Repeat



Parallel Approach Breakdown

- Set a step value equal to 1
- Loop from step value to number of processors
 - if(Process Id % (step * 2)):
 - Receive a recvsize from its partner
 - Allocate memory of recvsize
 - Receive a sorted array from the it's partner; Store it in the just allocated memory
 - Call mergeSubset with two arrays [process local array] and [the array just recvd]
 - Allocate memory of recvsize * 2 to store the array returned by mergeSubset
 - else:
 - Send your size to it's partner
 - Send your locally sorted array to it's partner
 - Increment step by step*2
- Break out of the loop

Merge Subsets

- Input: Two Sorted Subsets, Size of Subset
- Output: One Sorted Set
- Compare Elements in Subsets

Set1 { 1, 7, 11, 12, 15, 34, 35, 36 }

Set2 { 2, 3, 4, 10, 13, 20, 21, 22 }

Output { 1, 2, 3, 4, 7, 10, 11, 12, 13, 15, 20, 21, 22, 34, 35, 36 }


Data Breakdown

- Both Serial Parallel programs read from a file of random numbers.

IMPORTANT NOTE: Since they read from a file they read the EXACT same set of numbers this is extremely important in determining the relative speed among the Serial and Parallel algorithms

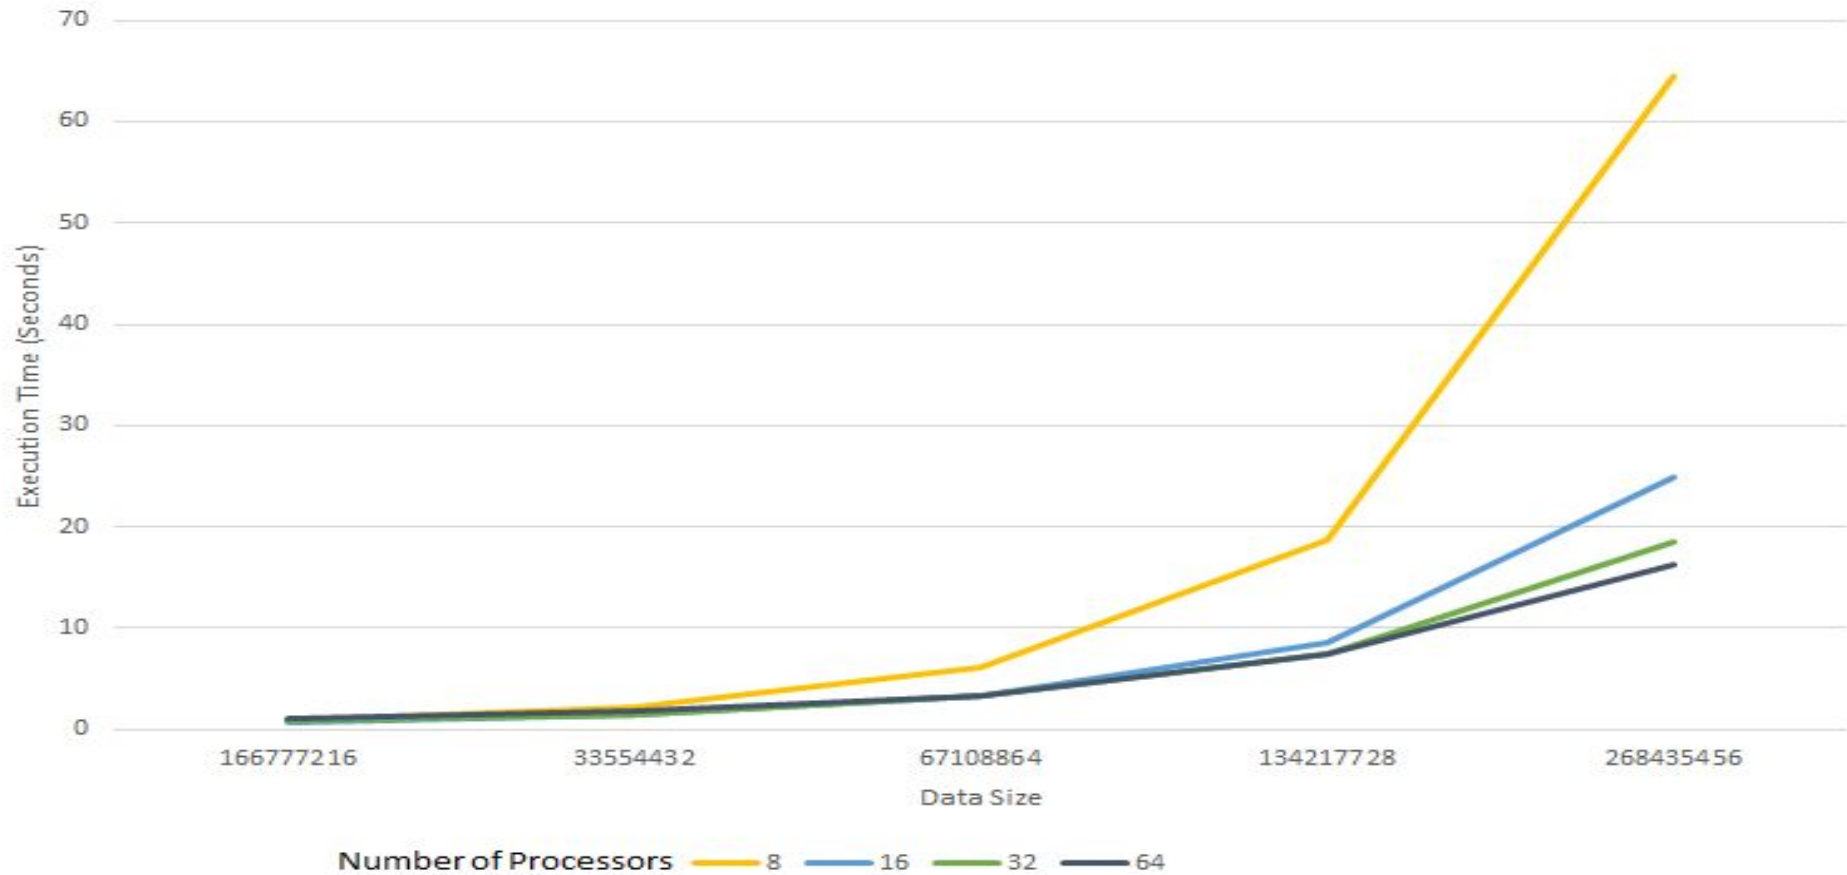
- Each of these files contain 2^x random numbers
 - In this presentation we are using data sets of size
 - 16,777,216 (2^{24})
 - 33,554,432 (2^{25})
 - 67,108,854 (2^{26})
 - 134,217,728 (2^{27})
 - 268,435,456 (2^{28})

Execution Time



Note: To sequentially sort the data size 2^{28} it takes 2989 seconds or 49 minutes and 48 seconds

Execution Time



Note: Where it previously took 49 minutes and 48 seconds for 2^{28} numbers to be sorted sequentially
Utilizing multiple processors can sort in a fraction of a time - FOR EXAMPLE: it takes 16 processors 24.88 seconds

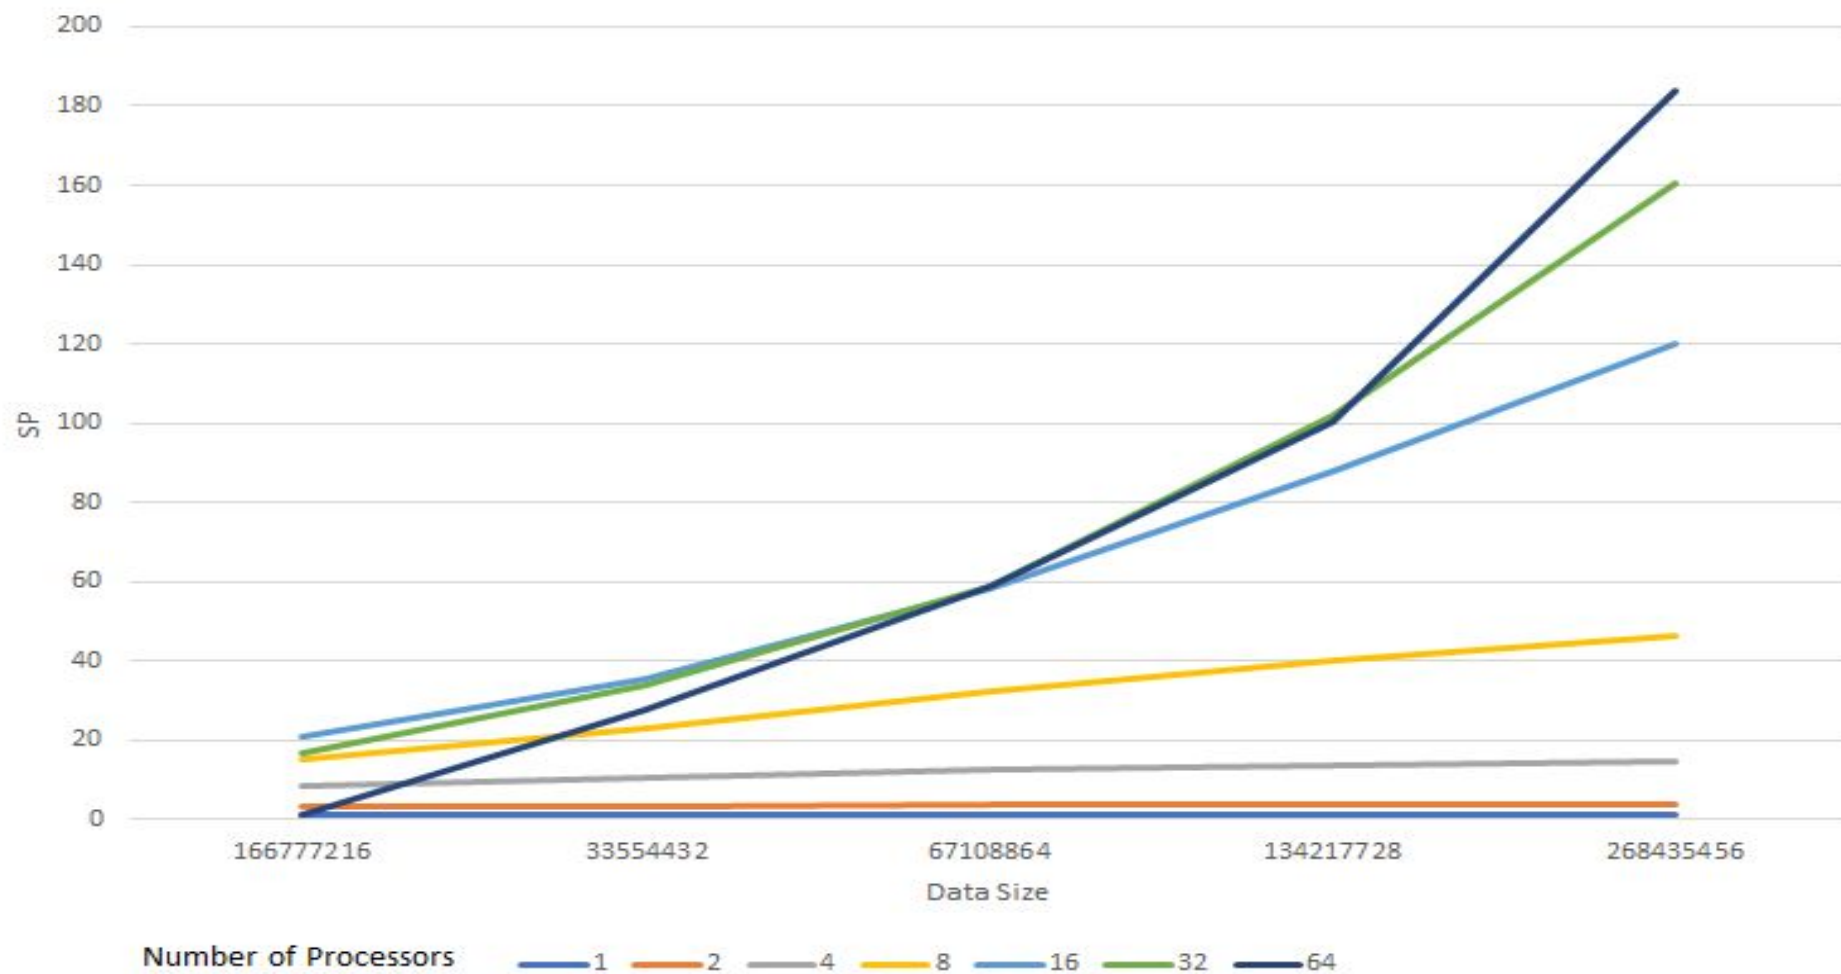
What is Speedup?

Speedup provides an answer to the question of how much faster can our multiprocessor implementation sort the given data set

$$\text{Speedup} = T_s / T_p$$

- T_s : Execution time of the serial approach
- T_p : Execution time of the parallel approach

Speedup



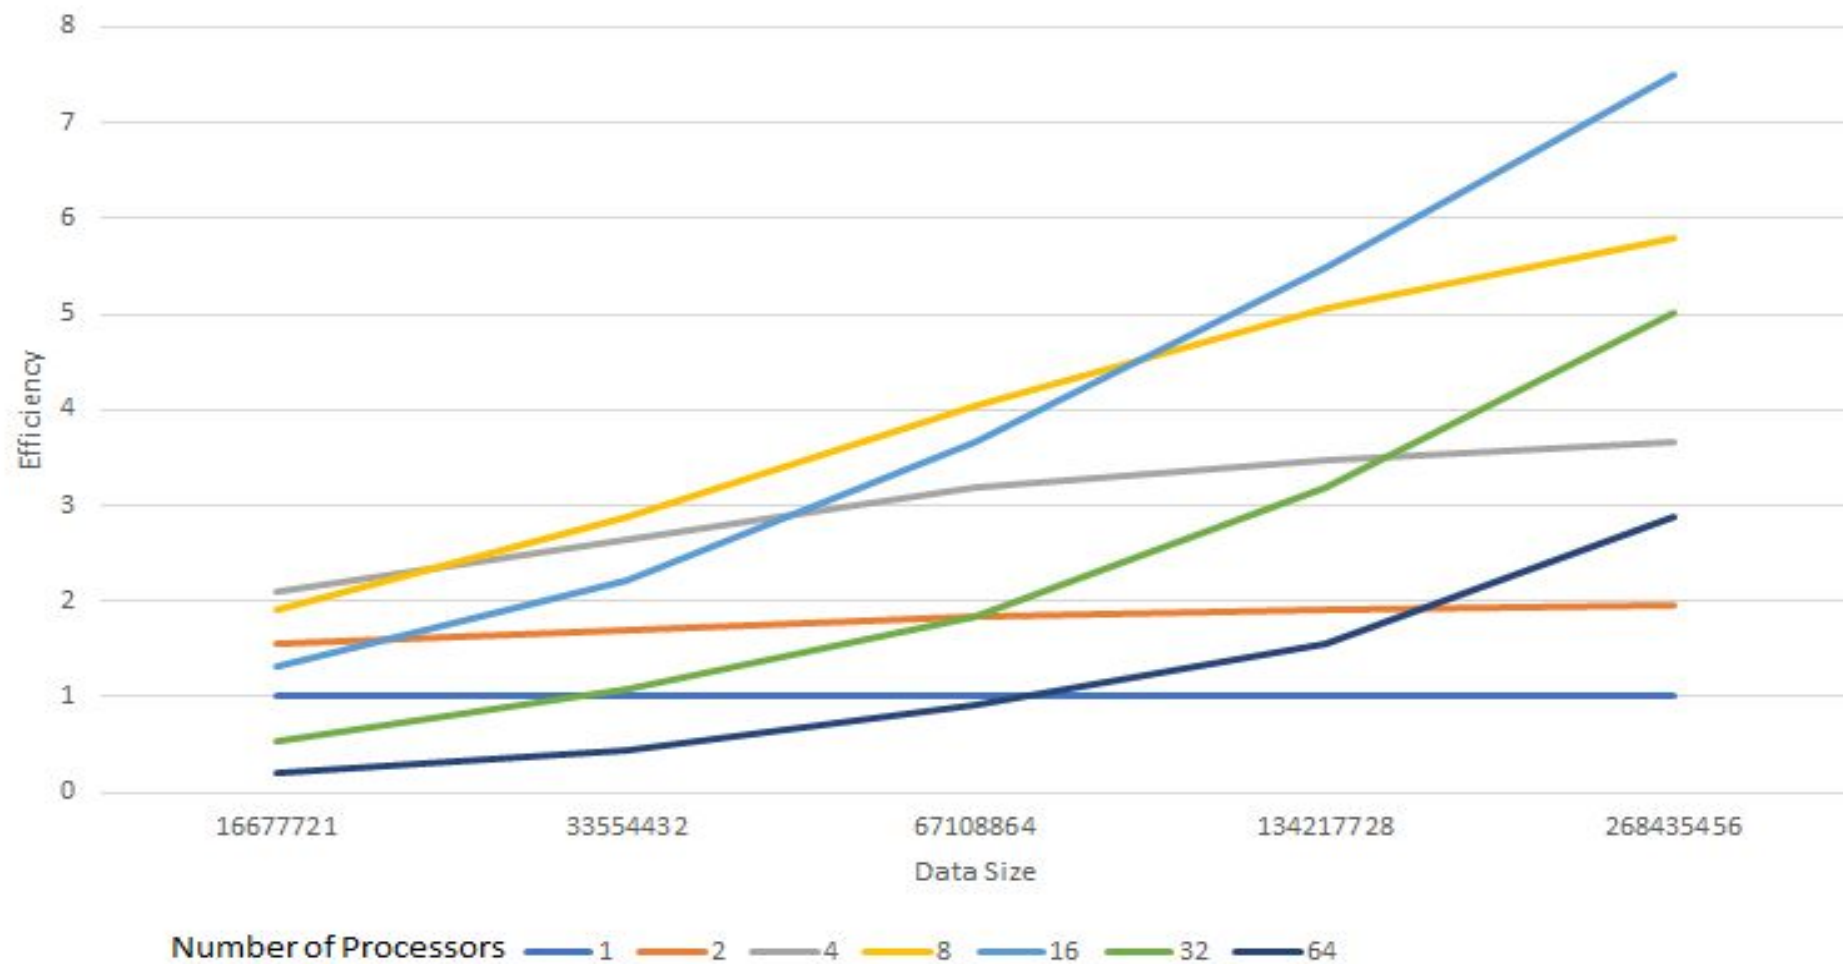
What is Efficiency?

Efficiency tells us how long each individual processor is being used on computation

Efficiency = $T_s / (\text{Number of Processors} * T_p)$

- T_s : Execution time of the serial approach
- T_p : Execution time of the parallel approach

Efficiency



Review

- Why do we need a parallel sort?
 - Large data sets
- Divide and conquer datasets
- Results
 - Execution times
 - Speed up
 - Efficiency

Thank You! Any Questions?

